


Neural Abstract Interpretation

Shaurya Gumber (sgomber2@illinois.edu)

Department of Computer Science, University of Illinois Urbana-Champaign, IL, USA

ABSTRACT INTERPRETATION

Helps analyze programs, neural networks, and real-world systems by interpreting them in an abstract domain!



ABSTRACT TRANSFORMERS

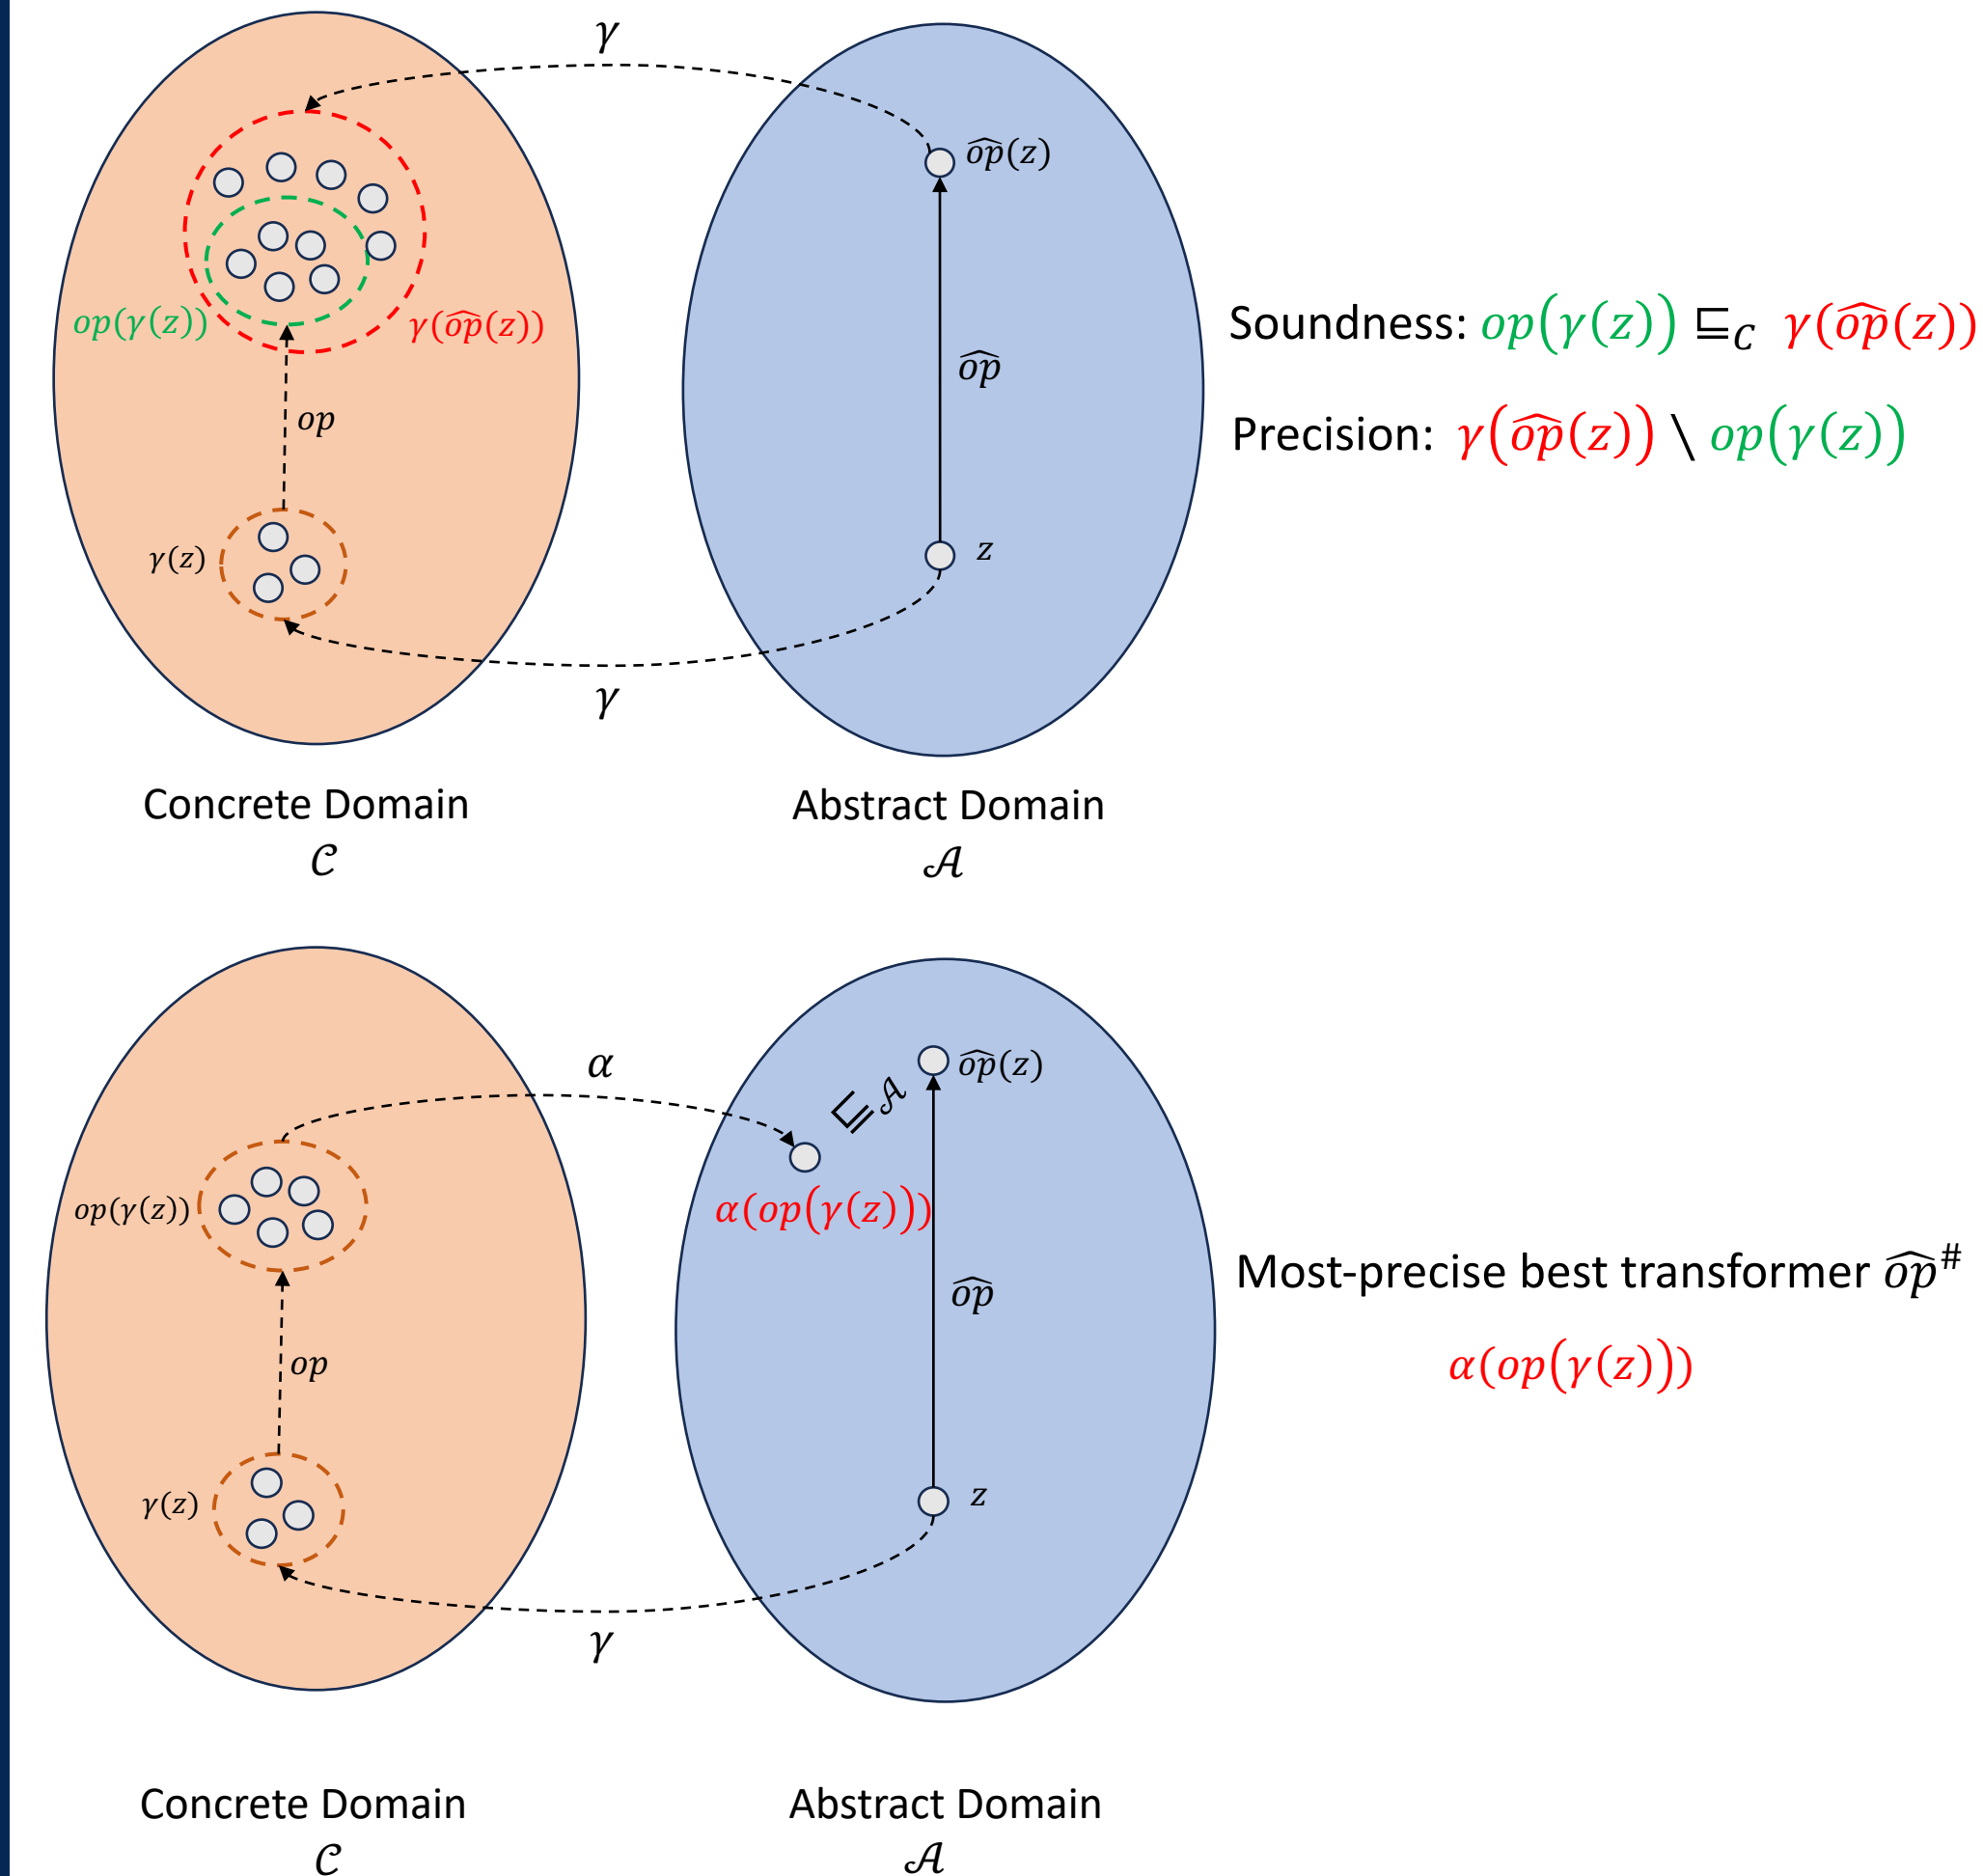
Operator op in $\mathcal{C} \rightarrow$ Need an equivalent \widehat{op} in \mathcal{A}

$x + y \rightarrow [1, 2] \hat{+} [3, 4]$

Interval Add

$x = 2x + 3y$

Octagon Affine Assignment



TRANSFORMERS: TEDIOUS TO IMPLEMENT!!

- Ensure soundness always.
- Most-precise implementations: **computationally expensive**. (join in Octagon domain is cubic and exponential in Polyhedra)
- Scalable implementations need **intricate optimizations**. (like octagon and polyhedral decomposition)
- Hand-crafted transformers can be **imprecise**. (like affine assignment in the octagon domain)
- Trade-off between Ease of Implementation, Soundness, Precision, and Efficiency.

MOTIVATES THE NEED TO LEARN SOUND AND PRECISE TRANSFORMERS AUTOMATICALLY!

GENERAL LEARNING PROBLEM IS HARD

Find the sound and most-precise abstract transformer \widehat{op} for op from a set of functions \mathcal{F}

$$\widehat{op} = \min_{f \in \mathcal{F}} \sum_{a \in \mathcal{A}} \mathcal{L}_P(\widehat{op}^\#(a), f(a)) \text{ s.t. } \sum_{a \in \mathcal{A}} \mathcal{L}_S(op, a, f(a)) = 0$$

PRECISION

- $\mathcal{L}_P(a_1, a_2)$ measures how “big” a_2 is as compared to a_1 .
- Ensures that learned transformer is closest to $\widehat{op}^\#$ in precision.

SOUNDNESS

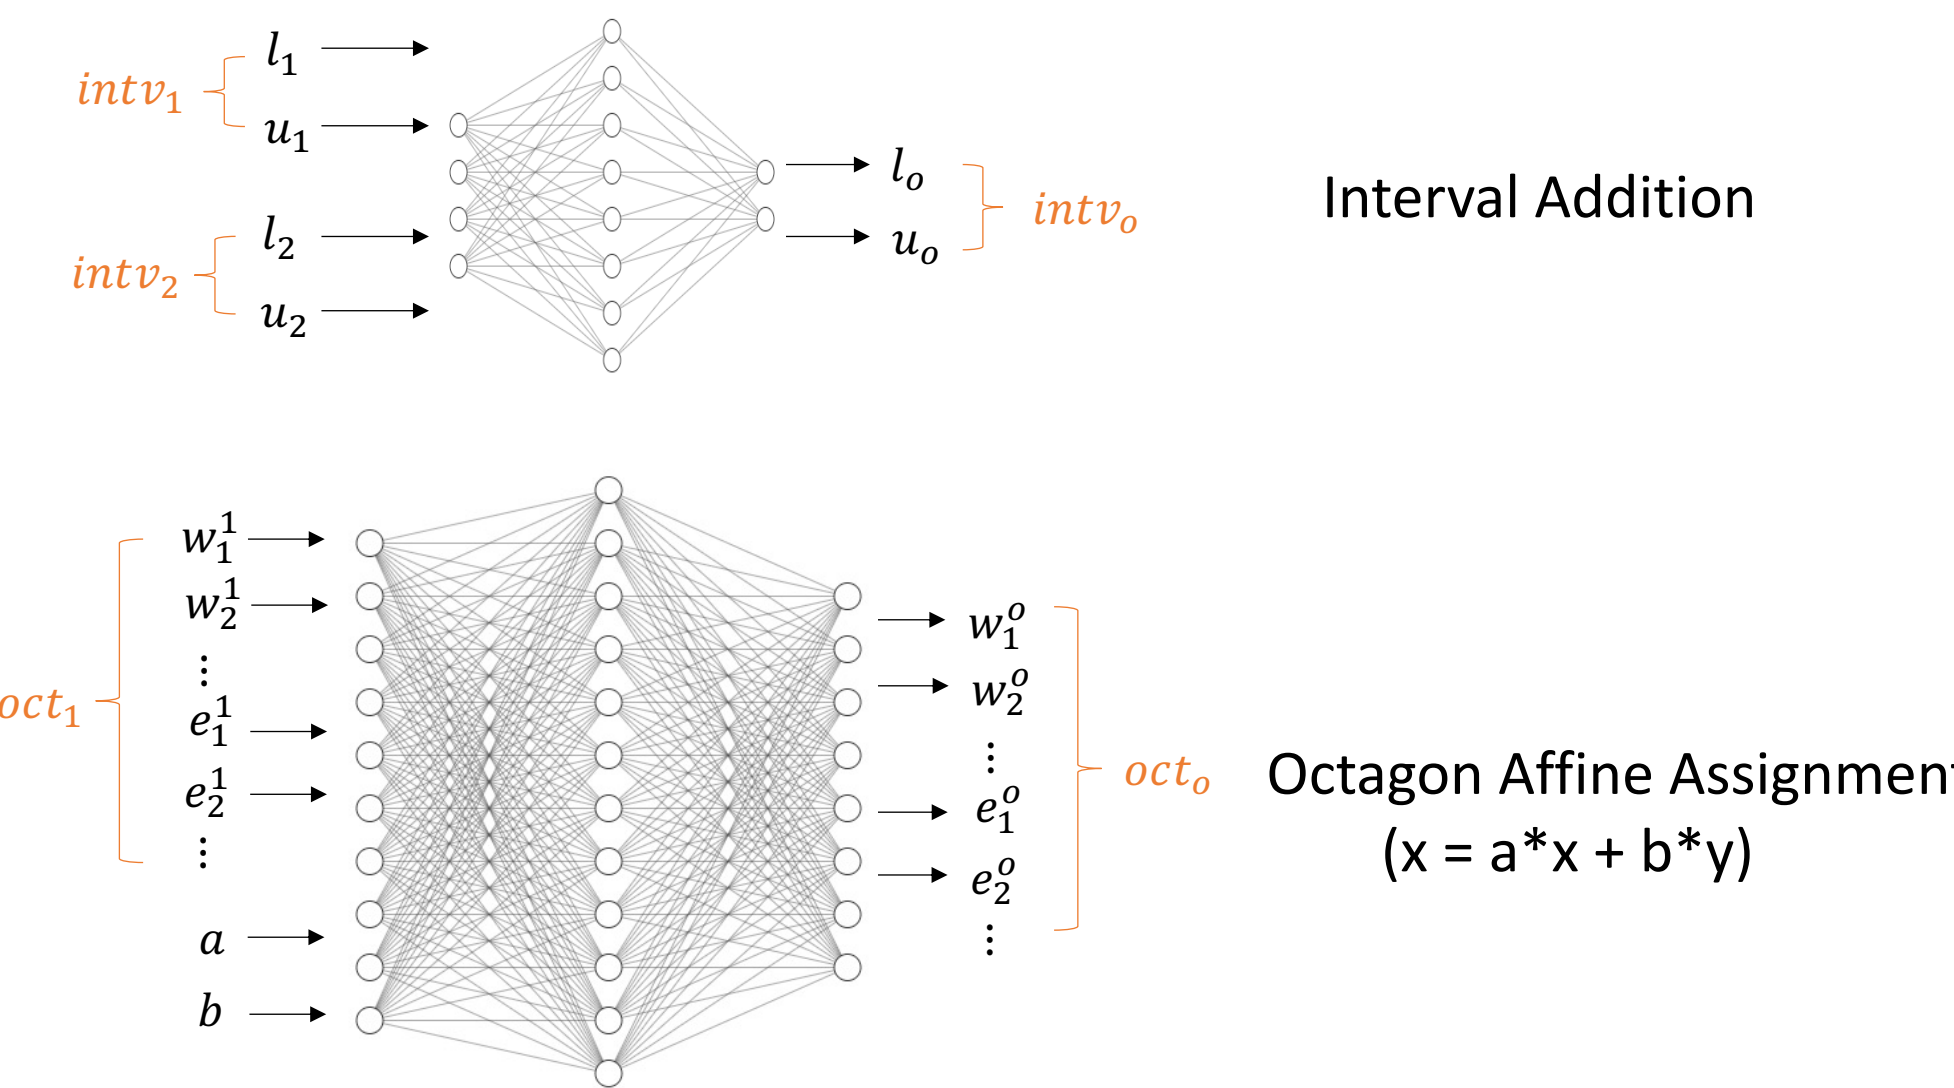
- $\mathcal{L}_S(op, a, f(a)) = 0$ if $f(a)$ is a sound output of f on a .
- Ensures f is sound on all $a \in \mathcal{A}$.

- \mathcal{A} has an infinite size!
- $\widehat{op}^\#$ just a specification $\rightarrow \mathcal{L}_P(a_1, a_2)$ thus hard to compute!
- $\mathcal{L}_S(op, a, f(a))$ checks for soundness: $op(\gamma(a)) \sqsubseteq_C \gamma(f(a))$
 \rightarrow via SMT
 \rightarrow expensive and not differentiable!

NEURAL ABSTRACT TRANSFORMERS

As neural networks effectively approximate complex functions, we propose:

NEURAL NETWORKS THAT SERVE AS ABSTRACT TRANSFORMERS!



Neural transformers allow for data-driven learning of transformers.

We propose **supervised** and **unsupervised** relaxations of the learning problem to train these transformers.

KEY BENEFITS OF NEURAL TRANSFORMERS

- Automatic generation of transformers with varying soundness and precision **eases development costs**.
- Neural transformers are **differentiable**, which allows us to pose and solve interesting problems like invariant generation as gradient-guided learning methods.
- Neural transformers can offer **faster alternatives** to computationally intensive transformers, such as the octagon join transformer, and **more precise alternatives** to transformers, like octagon affine assignment. (Unsound cases can be handled by resorting to hand-crafted transformers' outputs)

SUPERVISED LEARNING OF TRANSFORMERS

$$\min_{\theta} \mathbb{E}_{(X_i, Y_i) \sim \mathcal{D}} [\alpha * \mathcal{L}'_S(Y_i, N_{\theta}(X_i)) + \beta * \mathcal{L}'_P(Y_i, N_{\theta}(X_i))]$$

Soundness Enforcing Loss \mathcal{L}'_S

$\mathcal{L}'_S(a_1, a_2)$: Differentiable proxy for size of the set $\gamma(a_1) \setminus \gamma(a_2)$
 $\mathcal{L}'_S(a_1, a_2) = 0$ implies a_2 over-approximates a_1 : $\gamma(a_1) \sqsubseteq_C \gamma(a_2)$

Intervals: $\mathcal{L}'_S([l_1, u_1], [l_2, u_2]) = \max(l_2 - l_1, 0) + \max(u_1 - u_2, 0)$

$$\mathcal{L}'_S([1, 2], [0, 5]) = 0 \quad \mathcal{L}'_S([1, 6], [2, 4]) = (2-1) + (6-4) = 3 \quad (\text{Guides model's output to include } [1, 6])$$

Octagons: $\mathcal{L}'_S(o_1, o_2) = \sum_{i,j} \text{ite}(c_{ij} - c'_{ij} > 0, c_{ij} - c'_{ij}, 0)$ $o_1: \pm v_i \pm v_j \leq c_{ij}$
 $o_2: \pm v_i \pm v_j \leq c'_{ij}$

$$o_1: v_1 - v_2 \leq 5 \quad o_2: v_1 - v_2 \leq 3 \quad \mathcal{L}'_S(o_1, o_2) = 5 - 3 = 2$$

Precision Enforcing Loss \mathcal{L}'_P

$\mathcal{L}'_P(a_1, a_2)$: Differentiable approximation of how “big” a_2 is as compared to a_1 .
 Reducing it leads to “smaller” more-precise outputs.

Intervals: $\mathcal{L}'_P([l_1, u_1], [l_2, u_2]) = (u_2 - l_2) - (u_1 - l_1)$

$$\mathcal{L}'_P([1, 2], [0, 5]) = 5 - 1 = 4 \quad (\text{Guides model's output to match size of } [1, 2])$$

Octagons: $\mathcal{L}'_P(o_1, o_2) = \sum_{i,j} |c_{ij} - c'_{ij}|$ $o_1: \pm v_i \pm v_j \leq c_{ij}$
 $o_2: \pm v_i \pm v_j \leq c'_{ij}$

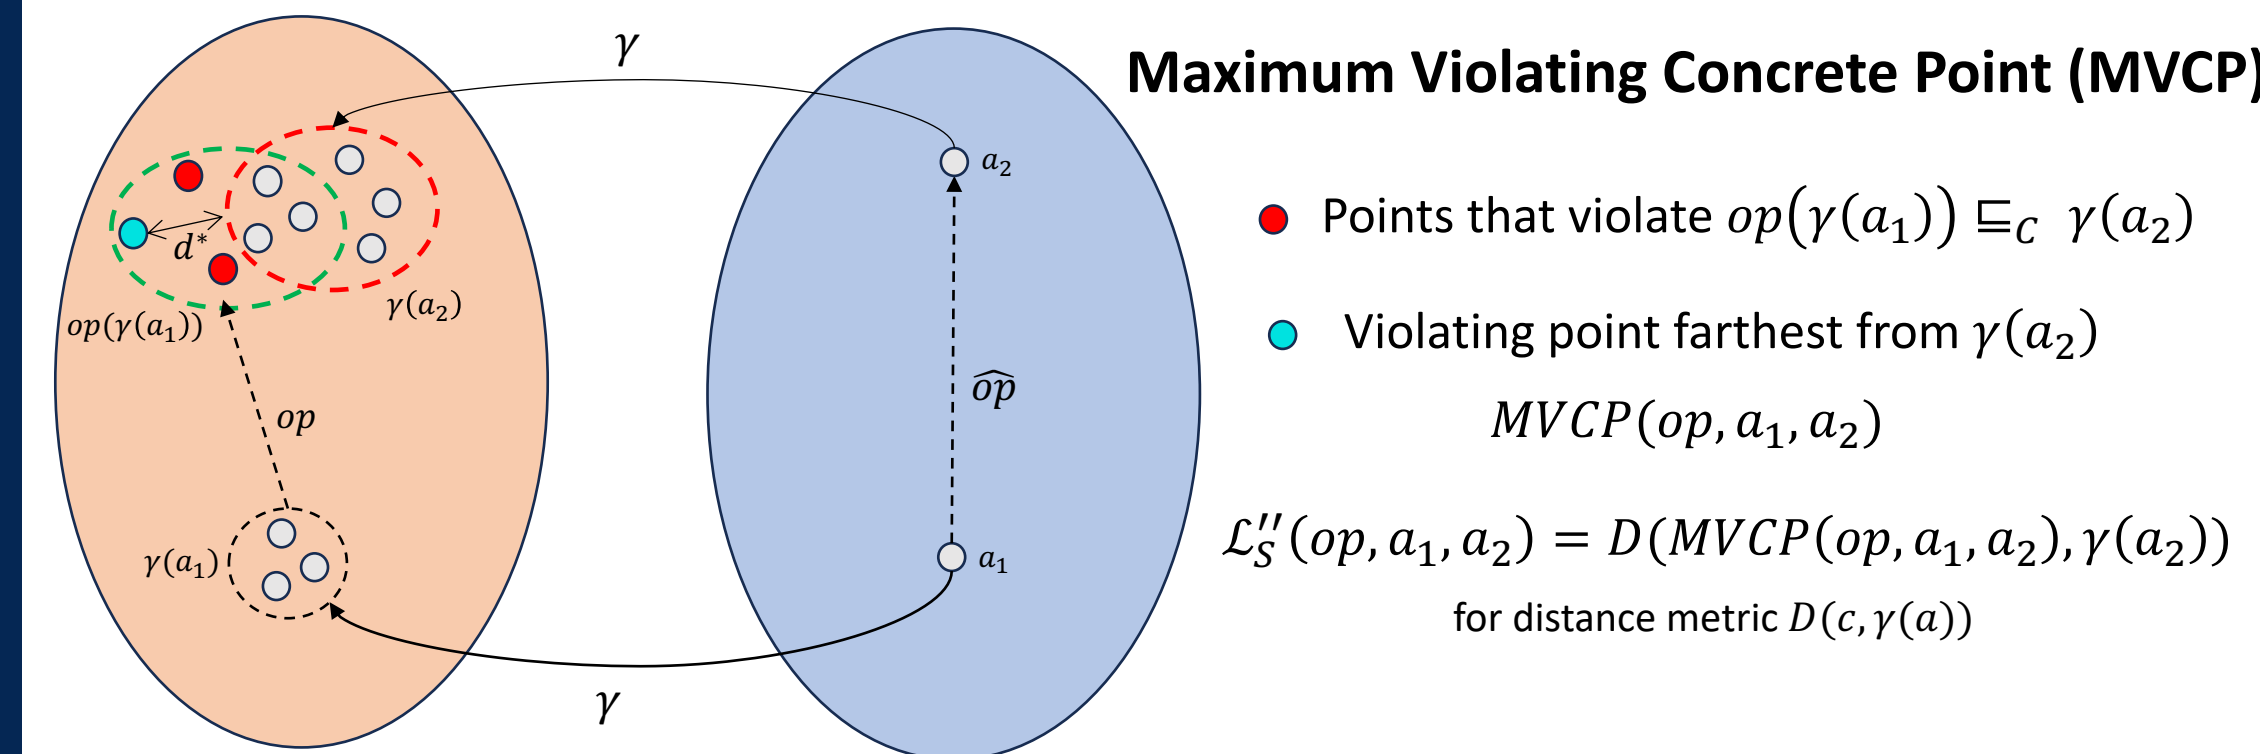
α and β are soundness and precision weights that manage the trade-off!

UNSUPERVISED LEARNING OF TRANSFORMERS

$$\min_{\theta} \mathbb{E}_{(X_i) \sim \mathcal{D}} [\alpha * \mathcal{L}''_S(op, X_i, N_{\theta}(X_i)) + \beta * \mathcal{L}_P''(N_{\theta}(X_i))]$$

Soundness Enforcing Loss \mathcal{L}''_S

$\mathcal{L}''_S(op, a_1, a_2)$ should guide a_2 to be a sound output of op on a_1 : $op(\gamma(a_1)) \sqsubseteq_C \gamma(a_2)$



Intervals


$$op = abs$$

$$a_1 = [-10, 15]$$

$$a_2 = [0, 12]$$

$$op(a_1) = [0, 15] \quad \mathcal{L}''_S = 3$$

Octagons



Precision Enforcing Loss \mathcal{L}''_P

$\mathcal{L}''_P(a)$ gives a differentiable approximation of the size of a .

Intervals

$$\mathcal{L}''_P([l, u]) = (u - l)$$

Octagons

$$\mathcal{L}''_P(o) = \sum_{i,j} |c_{ij}| \quad o: \pm v_i \pm v_j \leq c_{ij}$$

RESULTS

Supervised Learning

Weights (α, β)	Interval Abs		Interval Join	
	(Soundness, Precision)	Soundness (%)	Imprecision	Soundness (%)
(-, -)		20.03	4.44	3.88
(1, 1)		26.39	1.57	32.34
(2, 1)		47.43	5.74	40.53
(5, 1)		66.88	11.70	63.10
(7, 1)		84.02	10.39	73.07
(10, 1)		97.72	18.41	89.24
(50, 1)		99.99	40.63	99.57

Soundness %: Percentage of sound outputs on a test set of 10,000 input-output pairs.

Imprecision: The avg. difference in sizes of intervals produced by the model and the ground truth (for SOUND cases).

Soundness Weight (α)	Precision Weight (β)	Soundness Measure (%)	Imprecision Measure
1	1	9.6	49.96
10	1	36.6	85.30
20	1	49.0	110.51
50	1	64.5	129.26
100	1	79.0	184.58
250	1	86.3	291.23

Neural Octagon Join (3 variables)

Unsupervised Learning

Weights (α, β)	Interval Abs		Interval Join	
	(Soundness, Precision)	Soundness (%)	Imprecision	Soundness (%)
(-, -)		20.03	4.44	3.91
(20, 10)		25.04	4.29	38.99
(30, 10)		63.04	25.86	53.65
(50, 10)		85.96	36.95	93.03
(75, 10)		100	73.17	97.95

Soundness Weight (α)	Precision Weight (β)	Soundness Measure (%)	Imprecision Measure
10	10	20.6	85.30
20	10	35.1	150.51
50	10	57.1	229.26
100	10	80.0	384.58

Octagon affine assignment $x = ax + by$

Benefits of Differentiability: Invariant Search

Consider a while program: $\mathcal{P} = \langle \text{init} \rangle \text{ while}(B) \text{ do } C \text{ od}$

Find O_{inv} such that:
 $(O_{init} \in O_{inv}) \wedge (\hat{C}(O_{inv} \hat{\cap} B) \in O_{inv})$

Training neural transformers \hat{C}^* and \hat{B}^* for C and B

$$\mathcal{L}'_S(O_{init}, O_{inv}) + \mathcal{L}'_S(\hat{C}^*(O_{inv} \hat{\cap} B), O_{inv})$$

$$\mathcal{L}'_S(o_1, o_2) = \sum_{i,j} \text{ite}(c_{ij} - c'_{ij} > 0, c_{ij} - c'_{ij}, 0)$$

```
x = 100;
y = 150;
while (y <= 600)
{
  x = x + y;
  y = 2*y;
}
```

- $\{y \geq 65.51, x - y \leq -49.95, -x - y \leq 74.89\}$
- $\{x - y \leq 13.239\}$